Wikidata Human Gender Index

First Author
University of Author
Authortown, CA 94022, USA
author1@anotherco.edu

Second Author
VP, Authoring
Authorship Holdings, Ltd.
Awdur SA22 8PP, UK
author2@author.ac.uk

Fifth Author
YetAuthorCo, Inc.
Authortown, BC V6M 22P
Canada
author5@anotherco.com

Sixth Author
Université de Auteur-Sud
40222 Auteur France
author6@author.fr

Abstract
UPDATED—January 8, 2016. This sample paper describes the formatting requirements for SIGCHI Extended Abstract Format, and this sample file offers recommendations on writing for the worldwide SIGCHI readership. Please review this document even if you have submitted to SIGCHI conferences before, as some format details have changed relative to previous years. Abstracts should be about 150 words. Required.

Author Keywords
Authors' choice; of terms; separated; by semicolons; include commas, within terms only; required.

ACM Classification Keywords
H.5.m [Information interfaces and presentation (e.g., HCI)]: Miscellaneous; See [http://acm.org/about/class/1998/]: for full list of ACM classifiers. This section is required.

Introduction
Problematize introduction and make obvious need.

Landscape of ways in which Wikipedia shows bias, and Landscape of biases.

Our first observation is from September 17th 2014, and latest is January 3rd 2016. Although our dataset is now updated weekly following the official data dumps, automation

was not completed until June 28th 2015, and so there is a window missing from October 2014 to June 2015.

Albeit with caveats and huge biases, if the data is to believed, affords us an unprecedented look at gender representation on a time and geographic scale not possible before.

Introduce the dataset, and it's potential applications.

Dataset Details

The Data is collected by (I'm sure I've written this somewhere already).

How to use the data.

Where is it located. What types of files. How is it updated. How can they be used.

Show demonstrations of graphs from the website.

There is little cleaning done. More on this later.

Longitudinal Data

Total humans in Wikidata increased from 5,869,606 to 6,999,542, and shows linear, unconstrained growth.

We should also be curious to the data quality of the increasing humans. One way to think about this is about the how much data is accompanying these human entries. We looked at the properties citizenship, place of birth, and ethnic group which will help us best geographically place a human. Another mark of quality is whether a human in Wikidata has an entry in a Wikipedia - a "sitelink" in Wikidata vocabulary. **??** shows the rate of accompanying properties, at the earliest and latest snapshots from 2014 and 2016 respectively. The statistics show that data quality has been increasing uniformly over time. The number of humans with

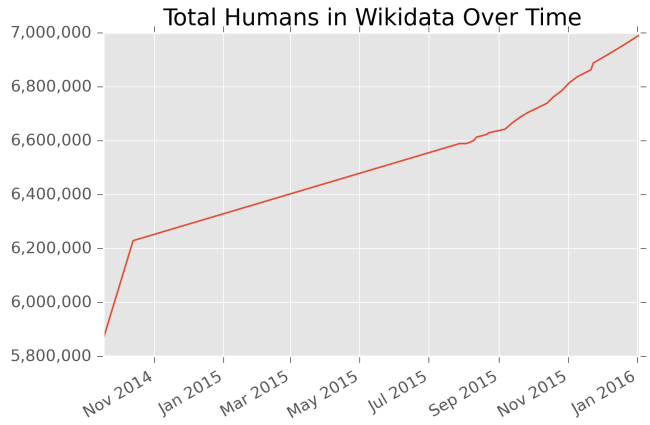

Figure 1: Total number of humans found in Wikidata at each snapshot period.

Table 1: Change in rates of human-accompanying properties

snapshot	gender	citizenship	place of birth	ethnic group	at least 1 s
2014-09-17	95.29%	42.82%	24.01%	0.31%	9
2016-01-03	96.54%	58.22%	30.51%	0.56%	9

gender data increased by over 1%, closer to complete coverage. Likewise *citizenship* data increased by 15% , *place of birth* by 6% , and *ethnic group* almost doubled 1. Curiously the rate of humans having sitelinks decreased slightly.

A Wikidata human without a Wikipedia article is know as a "structural item"; for instance a Monarch without a Wikipedia article, but is a needed link in a family tree. With the view that a structural item is an artefact from humans paying attention to Wikidata's structure, the decrease in sitelinked humans can also be seen as a rise in data quality.

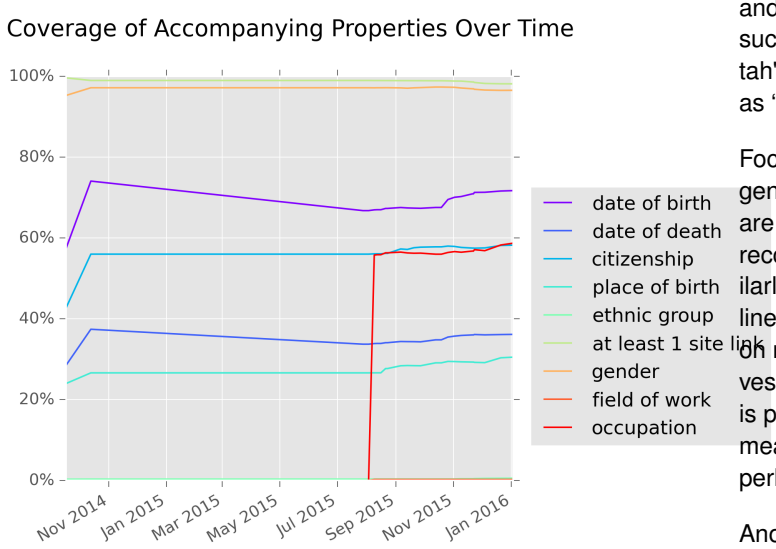

Figure 2: Trend of human-accompanying properties by snapshot.

Another important factor to note is that our dataset tries not to clean itself to fit any model. In fact the “gender" property in Wikidata is actually labelled in English “sex or gender" (no distinction), and not limited to any value. Over our time recording we found 36 values used for “sex or gender", including “male" and “female", but extending to nonbinary genders “transgender female", “intersex", “fa’afafine", “transgender", “Gender fluid", “genderqueer", “kathoey", and “queer". At times the other information is recorded such as “gay", or “homosexuality" and “Alien" or "cheetah". And even what seem to be mistakes are left in such as “Solanum tuberosum", “Messi", or “sociologist".

Focusing on one of our motivations, monitoring the trend of gender representation, we inspect the rate at which women are recorded in Wikidata. **??** shows the ratio of “female" recorded humans versus all gendered biographies. Similarly to total biographies this measure is rising at a fairly linear rate of approximately 0.5% per year. The final months on record show a slight decline which warrants further investigation. In fact being able to measure at a level where is precisely one of the points of having such a live-updating measure - to be able to detect trends as they happen, and perhaps relate them to community issues.

Another way in which Wikimedians can use the data is to look at trends specific to Wikipedia languages. It is easy to use the data to compile a top-10 list of Wikipedias whose female ratio of humans increased the most, see 2. This measure does not explain whether women’s representation in those languages increased because because editors took longer to record women’s gender in Wikidata and were catching up in the observed period, or that these languages became more women-focused over the snapshotting period.

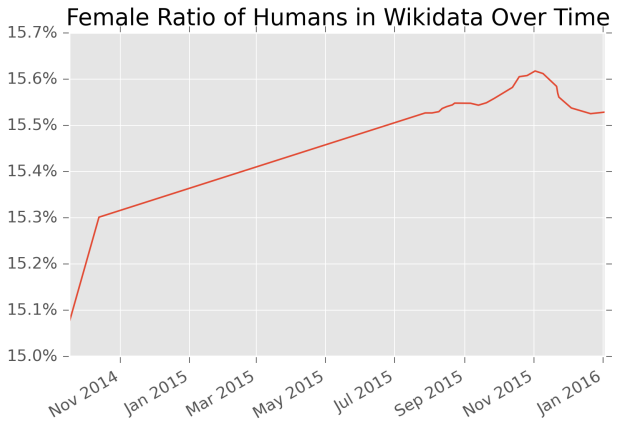

Figure 3: Trend of human-accompanying properties by snapshot.

Table 2: Top 10 Wikis by increase in female ratio of humans from October 2014 to January 2016

Wiki	Increase in female ratio of humans
Lithuanian	5.31%
Japanese	4.76%
Estonian	4.58%
Slovenian	2.19%
Tagalog	1.63%
Korean	1.38%
Finnish	1.33%
Wikimedia Commons	1.20%
Farsi	1.17%
Hebrew	1.17%

Table 3: WHGI-country correlation to external indices. Correlation is the Spearman ρ, and signficances are *p<0.05, **p<0.01

snapshot	GEI	SIGI	GGGI	GDI
2014-09-17	0.417**	0.338**	0.310*	0.278**
2016-01-03	0.457**	0.402**	0.386**	0.299**

Validation

We validated out data by comparing it against 3 exogenous measures. Wikidata date of birth frequency versus historical world population trends, Wikidata gender by country versus external by-country gender-disparity indexes, and Wikidata occupation gender versus United States Bureau of Labour Statistics occupation gender.

World Population

External Indices

Calibrated start dates were each 1900 or 1910 - a good sign. Over time the the WHGI-county is become more correlated to major external gender-disparity indexes 3. In context with the fact that data quality is rising, this can be taken to mean that as Wikidata becomes more complete it is modelling the real world more.

Potential Application

Open Knowledge Foundation founder Rufus Pollock once said “The best thing to do with your data will be thought of by someone else.âĂİ

FRB/wikisize. Is it data completeness or feminist-focus?

Use for determining women’s notability for historical events.

Linking with VIAF.

Comparing language’s inherent gender bias to shown.

Wikimedian communities can use this an introspective watch.

Useful in all the same ways that external indices like the UNDP are too.

References